

Robots

Video #1: Robots (Crash Course Computer Science #37)

Today we're going to talk about robots! Robots are often thought of as a technology of the future, but they're already here by the millions in the workplace, our homes, and pretty soon on the roads. We'll discuss the origins of robotics to its proliferation, and even look at some common control designs that were implemented to make them more useful in the workplace. Robots are often thought of as a menace or danger to society, and although there definitely is the propensity for malicious uses, robots also have the potential to drastically improve the world.

Video #1 Transcription

Hi, I'm Carrie Anne, and welcome to CrashCourse Computer Science! Today we're going to talk about robots.

The word "robot" was first used in a 1920 Czech play to denote artificial, humanoid characters. The word was derived from "robota", the slavic-language word for a forced laborer, indicating peasants in compulsory service in feudal, nineteenth century Europe. But even a century later, it's still a common portrayal: mass-produced, efficient, tireless creatures that look human-esque, but are emotionless, indifferent to self-preservation and lack creativity.

There are many definitions for robots, but in general, these are machines capable of carrying out a series of actions automatically, guided by computer control. How they look isn't part of the equation – robots can be industrial arms that spray paint cars, drones that fly, snake-like medical robots that assist surgeons, as well as humanoid robotic assistants. Although the term "robot" is sometimes applied to interactive virtual characters, it's more appropriate to call these "bots", or even better, "agents." That's because the term "robot" carries a physical connotation, a machine that lives in and acts on the real world.

The more general idea of self-operating machines goes back even further than the 1920s. Many ancient inventors created mechanical devices that performed functions automatically, like keeping the time and striking bells on the hour. There are plenty of examples of automated animal and humanoid figures that would perform dances, sing songs, strike drums, and do other physical actions. These non-electrical and certainly non-electronic machines were called automatons. For instance, an early automaton created in 1739 by the Frenchman Jacques de Vaucanson was the Canard Digerateur or Digesting Duck, a machine in the shape of a duck that appeared to eat grain and then defecate.

The first machines controlled by computers emerged in the late 1940s. These Computer Numerical Control, or CNC, machines could run programs that instructed a machine to perform a series of operations. This level of control also

enabled the creation of new manufactured goods, like milling a complex propellor design out of a block of aluminum – something that was difficult to do using standard machine tools, and with tolerances too small to be done by hand. CNC machines were a huge boon to industry, not just due to increased capability and precision, but also in terms of reducing labor costs by automating human jobs – a topic we’ll revisit in a later episode. The first commercial deployment was a programmable industrial robot called the Unimate, sold to General Motors in 1960 to lift hot pieces of metal from a die casting machine and stack them. This was the start of the robotics industry. Soon, robots were stacking pallets, welding parts, painting cars and much more.

The first image that jumps to your mind is probably a humanoid robot, like we usually see in shows or movies. Sometimes they’re our friends and colleagues, but more often, they’re sinister, apathetic, and battle-hardened. We also tend to think of robots as a technology of the future. But the reality is: they’re already here – by the millions – and they’re our workmates, helping us to do things harder, better, faster, and stronger.

For simple motions – like a robotic gripper that moves back and forth on a track – a robot can be instructed to move to a particular position, and it’ll keep moving in that direction until the desired position is reached, at which point it’ll stop. This behavior can be achieved through a simple control loop. First, sense the robot position. Are we there yet? Nope. So keep moving. Now sense position again. Are we there yet? Nope

So keep moving.

Now sense position again.

Are we there yet? Nope, so keep moving.

Are we there yet? Yes!

So we can stop moving, and also please be quiet!

Because we’re trying to minimize the distance between the sensed position and the desired position, this control loop is, more specifically, a negative feedback loop. A negative feedback control loop has three key pieces.

There’s a sensor, that measures things in the real world, like water pressure, motor position, air temperature, or whatever you’re trying to control.

From this measurement, we calculate how far we are from where we want to be – the error.

The error is then interpreted by a controller, which decides how to instruct the system to minimize that error.

Then, the system acts on the world through pumps, motors, heating elements, and other physical actuators.

In tightly controlled environments, simple control loops, like this, work OK. But in many real world applications, things are a tad more complicated.

Imagine that our gripper is really heavy, and even when the control loop says to stop, momentum causes the gripper to overshoot the desired position.

That would cause the control loop to take over again, this time backing the gripper up.

A badly tuned control loop might overshoot and overshoot and overshoot, and maybe even wobble forever.

To make matters worse, in real world settings, there are typically external and variable forces acting on a robot, like friction, wind and items of different weight. To handle this gracefully, more sophisticated control logic is needed.

A widely used control-loop, feedback mechanism is a proportional–integral–derivative controller. That’s a bit of a mouthful, so people call them PID controllers. These used to be mechanical devices, but now it’s all done in software.

Let’s imagine a robot that delivers coffee. Its goal is to travel between customers at two meters per second, which has been determined to be the ideal speed that’s both safe and expedient.

Of course, the environment doesn’t always cooperate. Sometimes there’s wind, and sometimes there’s uphill and downhill and all sorts of things that affect the speed of the robot.

So, it’s going to have to increase and decrease power to its motors to maintain the desired speed. Using the robot’s speed sensor, we can keep track of its actual speed and plot that alongside its desired speed.

PID controllers calculate three values from this data. First is the proportional value, which is the difference between the desired value and the actual value at the most recent instant in time or the present. This is what our simpler control loop used before. The bigger the gap between actual and desired, the harder you’ll push towards your target. In other words, it’s proportional control.

Next, the integral value is computed, which is the sum of error over a window of time, like the last few seconds. This look back helps compensate for steady state errors, resulting from things like motoring up a long hill. If this value is large, it means proportional control is not enough, and we have to push harder still.

Finally, there’s the derivative value, which is the rate of change between the desired and actual values. This helps account for possible future error and is sometimes called “anticipatory control”. For example, if you are screaming in towards your goal too fast, you’ll need to ease up a little to prevent overshoot. These three values are summed together, with different relative weights, to produce a controller output that’s passed to the system.

PID controllers are everywhere, from the cruise control in your car to drones that automatically adjust their rotor speeds to maintain level flight, as well as more

exotic robots, like this one that balances on a ball to move around. Advanced robots often require many control loops running in parallel, working together, managing everything from robot balance to limb position. Control loops are responsible for getting robot attributes like location to desired values.

So, you may be wondering where these values come from. This is the responsibility of higher-level robot software, which plans and executes robot actions, like plotting a path around sensed obstacles, or breaking down physical tasks, like picking up a ball, into simple, sequential motions.

Using these techniques, robots have racked up some impressive achievements – they’ve been to the deepest depths of Earth’s oceans and roved around on Mars for over a decade. But interestingly, lots of problems that are trivial for many humans have turned out to be devilishly difficult for robots: like walking on two legs, opening a door, picking up objects without crushing them, putting on a t-shirt, or petting a dog. These are tasks you may be able to do without thinking, but a supercomputer-powered robot fails at spectacularly. These sorts of tasks are all active areas of robotics research.

Artificial intelligence techniques, which we discussed a few episodes ago, are perhaps the most promising avenue to overcome these challenges. For example, Google has been running an experiment with a series of robotic arms that spend their days moving miscellaneous objects from one box to another, learning from trial and error. After thousands of hours of practice, the robots had cut their error rate in half. Of course, unlike humans, robots can run twenty-four hours a day and practice with many arms at the same time. So, it may just be a matter of time until they become adept at grasping things. But, for the time being, toddlers can out-grasp them.

One of the biggest and most visible robotic breakthroughs in recent years has been self-driving, autonomous cars. If you think about it, cars don’t have too many system inputs – you can speed up or slow down, and you can steer left or right. The tough part is sensing lanes, reading signs, and anticipating and navigating traffic, pedestrians, bicyclists, and a whole host of obstacles. In addition to being studded with proximity sensors, these robotic vehicles heavily rely on Computer Vision algorithms, which we discussed in Episode 35.

We’re also seeing the emergence of very primitive androids – robots that look and act like humans. Arguably, we’re not close on either of those goals, as they tend to look pretty weird and act even weirder. At least we’ll always have Westworld.

But anyway, these remain a tantalizing goal for roboticists that combine many computer science topics we’ve touched on over the last few episodes, like artificial intelligence, computer vision, and natural language processing. As for why humans are so fascinated by creating artificial embodiments of ourselves... you’ll have to go to Crash Course Philosophy for that. And for the foreseeable future, realistic androids will continue to be the stuff of science fiction.

Militaries also have a great interest in robots – they’re not only replaceable

but can surpass humans in attributes like strength, endurance, attention, and accuracy. Bomb disposal robots and reconnaissance drones are fairly common today. But fully autonomous, armed-to-the-teeth robots are slowly appearing, like the Samsung SGR-A1 sentry gun deployed by South Korea. Robots with the intelligence and capability to take human lives are called lethal autonomous weapons. And they're widely considered a complex and thorny issue.

Without doubt, these systems could save soldiers lives by taking them off the battlefield and out of harm's way. It might even discourage war altogether. Though it's worth noting that people said the same thing about dynamite and nuclear weapons. On the flip side, we might be creating ruthlessly efficient killing machines that don't apply human judgment or compassion to complex situations. And the fog of war is about as complex and murky as they come. These robots would be taking orders and executing them as efficiently as they can, and sometimes human orders turn out to be really bad.

This debate is going to continue for a long time, and pundits on both sides will grow louder as robotic technology improves. It's also an old debate – the danger was obvious to science fiction writer Isaac Asimov, who introduced a fictional “Three Laws of Robotics” in his 1942 short story “Runaround”. And then, later he added a zeroth rule. In short, it's a code of conduct or moral compass for robots – guiding them to do no harm, especially to humans. It's pretty inadequate for practical application and it leaves plenty of room for equivocation. But still, Asimov's laws inspired a ton of science fiction and academic discussion, and today there are whole conferences on robot ethics.

Importantly, Asimov crafted his fictional rules as a way to push back on “Robot as a Menace” memes common in fiction from his childhood. These were stories where robots went off the rails, harming or even destroying their creators in the process. Asimov, on the other hand, envisioned robots as useful, reliable, and even lovable machines. And it's this duality I want to leave you thinking about today. Like many of the technologies we've discussed throughout this series, there are benevolent and malicious uses. Our job is to carefully reflect on computing's potential and peril, and wield our inventive talents to improve the state of the world. And robots are one of the most potent reminders of this responsibility. I'll see you next week.